

# SQL for Data Analysis Cheat Sheet

## SQL

**SQL**, or *Structured Query Language*, is a language for talking to databases. It lets you select specific data and build complex reports. Today, SQL is a universal language of data, used in practically all technologies that process data.

## SELECT

Fetch the id and name columns from the product table:  
`SELECT id, name  
FROM product;`

Concatenate the name and the description to fetch the full description of the products:  
`SELECT name || ' - ' || description  
FROM product;`

Fetch names of products with prices above 15:  
`SELECT name  
FROM product  
WHERE price > 15;`

Fetch names of products with prices between 50 and 150:  
`SELECT name  
FROM product  
WHERE price BETWEEN 50 AND 150;`

Fetch names of products that are not watches:  
`SELECT name  
FROM product  
WHERE name != 'watch';`

Fetch names of products that start with a 'P' or end with an 's':  
`SELECT name  
FROM product  
WHERE name LIKE 'P%' OR name LIKE '%s';`

Fetch names of products that start with any letter followed by 'rain' (like 'train' or 'grain'):  
`SELECT name  
FROM product  
WHERE name LIKE '_rain';`

Fetch names of products with non-null prices:  
`SELECT name  
FROM product  
WHERE price IS NOT NULL;`

## GROUP BY

| PRODUCT  |             |  |  |
|----------|-------------|--|--|
| name     | category    |  |  |
| Knife    | Kitchen     |  |  |
| Pot      | Kitchen     |  |  |
| Mixer    | Kitchen     |  |  |
| Jeans    | Clothing    |  |  |
| Sneakers | Clothing    |  |  |
| Leggings | Clothing    |  |  |
| Smart TV | Electronics |  |  |
| Laptop   | Electronics |  |  |

  

| category    | count |
|-------------|-------|
| Kitchen     | 3     |
| Clothing    | 3     |
| Electronics | 2     |

## AGGREGATE FUNCTIONS

Count the number of products:  
`SELECT COUNT(*)  
FROM product;`

Count the number of products with non-null prices:  
`SELECT COUNT(price)  
FROM product;`

Count the number of unique category values:  
`SELECT COUNT(DISTINCT category)  
FROM product;`

Get the lowest and the highest product price:  
`SELECT MIN(price), MAX(price)  
FROM product;`

Find the total price of products for each category:  
`SELECT category, SUM(price)  
FROM product  
GROUP BY category;`

Find the average price of products for each category whose average is above 3.0:  
`SELECT category, AVG(price)  
FROM product  
GROUP BY category  
HAVING AVG(price) > 3.0;`

## ORDER BY

Fetch product names sorted by the price column in the default ASCending order:  
`SELECT name  
FROM product  
ORDER BY price [ASC];`

Fetch product names sorted by the price column in DESCending order:  
`SELECT name  
FROM product  
ORDER BY price DESC;`

## COMPUTATIONS

Use +, -, \*, / to do basic math. To get the number of seconds in a week:  
`SELECT 60 * 60 * 24 * 7;  
-- result: 604800`

## ROUNDING NUMBERS

Round a number to its nearest integer:  
`SELECT ROUND(1234.56789);  
-- result: 1235`

Round a number to two decimal places:  
`SELECT ROUND(AVG(price), 2)  
FROM product  
WHERE category_id = 21;  
-- result: 124.56`

## TROUBLESHOOTING

### INTEGER DIVISION

In PostgreSQL and SQL Server, the / operator performs integer division for integer arguments. If you do not see the number of decimal places you expect, it is because you are dividing between two integers. Cast one to decimal:  
`123 / 2 -- result: 61  
CAST(123 AS decimal) / 2 -- result: 61.5`

### DIVISION BY 0

To avoid this error, make sure the denominator is not 0. You may use the NULLIF() function to replace 0 with a NULL, which results in a NULL for the entire expression:  
`count / NULLIF(count_all, 0)`

## JOIN

JOIN is used to fetch data from multiple tables. To get the names of products purchased in each order, use:

```
SELECT
  orders.order_date,
  product.name AS product,
  amount
FROM orders
JOIN product
  ON product.id = orders.product_id;
```

Learn more about JOINS in our interactive [SQL JOINS](#) course.

## INSERT

To insert data into a table, use the INSERT command:  
`INSERT INTO category  
VALUES  
(1, 'Home and Kitchen'),  
(2, 'Clothing and Apparel');`

You may specify the columns to which the data is added. The remaining columns are filled with predefined default values or NULLs.  
`INSERT INTO category (name)  
VALUES ('Electronics');`

## UPDATE

To update the data in a table, use the UPDATE command:  
`UPDATE category  
SET  
 is_active = true,  
 name = 'Office'  
WHERE name = 'Office';`

## DELETE

To delete data from a table, use the DELETE command:  
`DELETE FROM category  
WHERE name IS NULL;`

Check out our interactive course [How to INSERT, UPDATE, and DELETE Data in SQL](#).

# SQL for Data Analysis Cheat Sheet

## DATE AND TIME

There are 3 main time-related types: **date**, **time**, and **timestamp**. Time is expressed using a 24-hour clock, and it can be as vague as just hour and minutes (e.g., 15:30 – 3:30 p.m.) or as precise as microseconds and time zone (as shown below):

2021-12-31 14:39:53.662522-05

|           |      |
|-----------|------|
| date      | time |
| timestamp |      |

YYYY-mm-dd HH:MM:SS.ssssss±TZ

14:39:53.662522-05 is almost 2:40 p.m. CDT (e.g., in Chicago; in UTC it'd be 7:40 p.m.). The letters in the above example represent:

- |                                                                                                                                                                                       |                                                                                                                                                                                                                                                                                                                                                                                                                 |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>In the date part:</b>                                                                                                                                                              | <b>In the time part:</b>                                                                                                                                                                                                                                                                                                                                                                                        |
| <ul style="list-style-type: none"><li>• YYYY – the 4-digit year.</li><li>• mm – the zero-padded month (01—January through 12—December).</li><li>• dd – the zero-padded day.</li></ul> | <ul style="list-style-type: none"><li>• HH – the zero-padded hour in a 24-hour clock.</li><li>• MM – the minutes.</li><li>• SS – the seconds. <i>Omissible</i>.</li><li>• ssssss – the smaller parts of a second – they can be expressed using 1 to 6 digits. <i>Omissible</i>.</li><li>• ±TZ – the timezone. It must start with either + or –, and use two digits relative to UTC. <i>Omissible</i>.</li></ul> |

## CURRENT DATE AND TIME

Find out what time it is:  
**SELECT CURRENT\_TIME;**

Get today's date:  
**SELECT CURRENT\_DATE;**  
In SQL Server:  
**SELECT GETDATE();**

Get the timestamp with the current date and time:  
**SELECT CURRENT\_TIMESTAMP;**

## CREATING DATE AND TIME VALUES

To create a date, time, or timestamp, write the value as a string and cast it to the proper type.

```
SELECT CAST('2021-12-31' AS date);
SELECT CAST('15:31' AS time);
SELECT CAST('2021-12-31 23:59:29+02'
AS timestamp);
SELECT CAST('15:31.124769' AS time);
```

Be careful with the last example – it is interpreted as 15 minutes 31 seconds and 124769 microseconds! It is always a good idea to write 00 for hours explicitly: '00:15:31.124769'.

## SORTING CHRONOLOGICALLY

Using ORDER BY on date and time columns sorts rows chronologically from the oldest to the most recent:  
**SELECT order\_date, product, quantity**  
**FROM sales**  
**ORDER BY order\_date;**

| order_date | product  | quantity |
|------------|----------|----------|
| 2023-07-22 | Laptop   | 2        |
| 2023-07-23 | Mouse    | 3        |
| 2023-07-24 | Sneakers | 10       |
| 2023-07-24 | Jeans    | 3        |
| 2023-07-25 | Mixer    | 2        |

Use the DESCending order to sort from the most recent to the oldest:

```
SELECT order_date, product, quantity
FROM sales
ORDER BY order_date DESC;
```

## COMPARING DATE AND TIME VALUES

You may use the comparison operators <, <=, >, >=, and = to compare date and time values. Earlier dates are less than later ones. For example, 2023-07-05 is "less" than 2023-08-05.

Find sales made in July 2023:  
**SELECT order\_date, product\_name, quantity**  
**FROM sales**  
**WHERE order\_date >= '2023-07-01'**  
**AND order\_date < '2023-08-01';**

Find customers who registered in July 2023:  
**SELECT registration\_timestamp, email**  
**FROM customer**  
**WHERE registration\_timestamp >= '2023-07-01'**  
**AND registration\_timestamp < '2023-08-01';**

**Note:** Pay attention to the end date in the query. The upper bound '2023-08-01' is not included in the range. The timestamp '2023-08-01' is actually the timestamp '2023-08-01 00:00:00.0'. The comparison operator < is used to ensure the selection is made for all timestamps less than '2023-08-01 00:00:00.0', that is, all timestamps in July 2023, even those close to the midnight of August 1, 2023.

## INTERVALS

An interval measures the difference between two points in time. For example, the interval between 2023-07-04 and 2023-07-06 is 2 days.

To define an interval in SQL, use this syntax:  
**INTERVAL '1' DAY**

The syntax consists of three elements: the INTERVAL keyword, a quoted value, and a time part keyword. You may use the following time parts: YEAR, MONTH, DAY, HOUR, MINUTE, and SECOND.

## Adding intervals to date and time values

You may use + or – to add or subtract an interval to date or timestamp values.

Subtract one year from 2023-07-05:  
**SELECT CAST('2023-07-05' AS TIMESTAMP)**  
**– INTERVAL '1' year;**  
**-- result: 2022-07-05 00:00:00**

Find customers who placed the first order within a month from the registration date:  
**SELECT id**  
**FROM customers**  
**WHERE first\_order\_date >**  
**registration\_date + INTERVAL '1' month;**

## Filtering events to those in the last 7 days

To find the deliveries scheduled for the last 7 days, use:  
**SELECT delivery\_date, address**  
**FROM sales**  
**WHERE delivery\_date <= CURRENT\_DATE**  
**AND delivery\_date >= CURRENT\_DATE**  
**– INTERVAL '7' DAY;**

**Note:** In SQL Server, intervals are not implemented – use the DATEADD() and DATEDIFF() functions.

## Filtering events to those in the last 7 days in SQL Server

To find the sales made within the last 7 days, use:  
**SELECT delivery\_date, address**  
**FROM sales**  
**WHERE delivery\_date <= GETDATE()**  
**AND delivery\_date >=**  
**DATEADD(DAY, -7, GETDATE());**

## EXTRACTING PARTS OF DATES

The standard SQL syntax to get a part of a date is  
**SELECT EXTRACT(YEAR FROM order\_date)**  
**FROM sales;**

You may extract the following fields:  
YEAR, MONTH, DAY, HOUR, MINUTE, and SECOND.

The standard syntax does not work in SQL Server. Use the DATEPART(part, date) function instead.  
**SELECT DATEPART(YEAR, order\_date)**  
**FROM sales;**

## GROUPING BY YEAR AND MONTH

Find the count of sales by month:

```
SELECT
EXTRACT(YEAR FROM order_date) AS year,
EXTRACT(MONTH FROM order_date) AS month,
COUNT(*) AS count
FROM sales
GROUP BY
year,
month
ORDER BY
year
month;
```

| year | month | count |
|------|-------|-------|
| 2022 | 8     | 51    |
| 2022 | 9     | 58    |
| 2022 | 10    | 62    |
| 2022 | 11    | 76    |
| 2022 | 12    | 85    |
| 2023 | 1     | 71    |
| 2023 | 2     | 69    |

Note that you must group by both the year and the month. EXTRACT(MONTH FROM order\_date) only extracts the month number (1, 2, ..., 12). To distinguish between months from different years, you must also group by year.

More about working with date and time values in our interactive [Standard SQL Functions](#) course.



# SQL for Data Analysis Cheat Sheet

## CASE WHEN

CASE WHEN lets you pass conditions (as in the WHERE clause), evaluates them in order, then returns the value for the first condition met.

```
SELECT
  name,
  CASE
    WHEN price > 150 THEN 'Premium'
    WHEN price > 100 THEN 'Mid-range'
    ELSE 'Standard'
  END AS price_category
FROM product;
```

Here, all products with prices above 150 get the *Premium* label, those with prices above 100 (and below 150) get the *Mid-range* label, and the rest receives the *Standard* label.

## CASE WHEN and GROUP BY

You may combine CASE WHEN and GROUP BY to compute object statistics in the categories you define.

```
SELECT
  CASE
    WHEN price > 150 THEN 'Premium'
    WHEN price > 100 THEN 'Mid-range'
    ELSE 'Standard'
  END AS price_category,
  COUNT(*) AS products
FROM product
GROUP BY price_category;
```

Count the number of large orders for each customer using CASE WHEN and SUM():

```
SELECT
  customer_id,
  SUM(
    CASE WHEN quantity > 10
      THEN 1 ELSE 0 END
  ) AS large_orders
FROM sales
GROUP BY customer_id;
```

... or using CASE WHEN and COUNT():

```
SELECT
  customer_id,
  COUNT(
    CASE WHEN quantity > 10
      THEN order_id END
  ) AS large_orders
FROM sales
GROUP BY customer_id;
```

## GROUP BY EXTENSIONS

### GROUPING SETS

GROUPING SETS lets you specify multiple sets of columns to group by in one query.

```
SELECT region, product, COUNT(order_id)
FROM sales
GROUP BY
  GROUPING SETS ((region, product), ());
```

| region | product | count |
|--------|---------|-------|
| USA    | Laptop  | 10    |
| USA    | Mouse   | 5     |
| UK     | Laptop  | 6     |
| NULL   | NULL    | 21    |

GROUP BY (region, product)

GROUP BY () – all rows

### CUBE

CUBE generates groupings for all possible subsets of the GROUP BY columns.

```
SELECT region, product, COUNT(order_id)
FROM sales
GROUP BY CUBE (region, product);
```

| region | product | count |
|--------|---------|-------|
| USA    | Laptop  | 10    |
| USA    | Mouse   | 5     |
| UK     | Laptop  | 6     |
| USA    | NULL    | 15    |
| UK     | NULL    | 6     |
| NULL   | Laptop  | 16    |
| NULL   | Mouse   | 5     |
| NULL   | NULL    | 21    |

GROUP BY region, product

GROUP BY region

GROUP BY product

GROUP BY () – all rows

### ROLLUP

ROLLUP adds new levels of grouping for subtotals and grand totals.

```
SELECT region, product, COUNT(order_id)
FROM sales
GROUP BY ROLLUP (region, product);
```

| region | product | count |
|--------|---------|-------|
| USA    | Laptop  | 10    |
| USA    | Mouse   | 5     |
| UK     | Laptop  | 6     |
| USA    | NULL    | 15    |
| UK     | NULL    | 6     |
| NULL   | NULL    | 21    |

GROUP BY region, product

GROUP BY region

GROUP BY () – all rows

## COALESCE

COALESCE replaces the first NULL argument with a given value. It is often used to display labels with GROUP BY extensions.

```
SELECT region,
  COALESCE(product, 'All'),
  COUNT(order_id)
FROM sales
GROUP BY ROLLUP (region, product);
```

| region | product | count |
|--------|---------|-------|
| USA    | Laptop  | 10    |
| USA    | Mouse   | 5     |
| USA    | All     | 15    |
| UK     | Laptop  | 6     |
| UK     | All     | 6     |
| All    | All     | 21    |

## COMMON TABLE EXPRESSIONS

A common table expression (CTE) is a named temporary result set that can be referenced within a larger query. They are especially useful for complex aggregations and for breaking down large queries into more manageable parts.

```
WITH total_product_sales AS (
  SELECT product, SUM(profit) AS total_profit
  FROM sales
  GROUP BY product
)
```

```
SELECT AVG(total_profit)
FROM total_product_sales;
```

Check out our hands-on courses on [Common Table Expressions](#) and [GROUP BY Extensions](#).

## WINDOW FUNCTIONS

Window functions compute their results based on a sliding window frame, a set of rows related to the current row. Unlike aggregate functions, window functions do not collapse rows.

### COMPUTING THE PERCENT OF TOTAL WITHIN A GROUP

```
SELECT product, brand, profit,
  (100.0 * profit /
    SUM(profit) OVER(PARTITION BY brand)) AS perc
FROM sales;
```

| product | brand  | profit | perc |
|---------|--------|--------|------|
| Knife   | Culina | 1000   | 25   |
| Pot     | Culina | 3000   | 75   |
| Doll    | Toyze  | 2000   | 40   |
| Car     | Toyze  | 3000   | 60   |

## RANKING

Rank products by price:

```
SELECT RANK() OVER(ORDER BY price), name
FROM product;
```

### RANKING FUNCTIONS

RANK – gives the same rank for tied values, leaves gaps.

DENSE\_RANK – gives the same rank for tied values without gaps.

ROW\_NUMBER – gives consecutive numbers without gaps.

| name     | rank | dense_rank | row_number |
|----------|------|------------|------------|
| Jeans    | 1    | 1          | 1          |
| Leggings | 2    | 2          | 2          |
| Leggings | 2    | 2          | 3          |
| Sneakers | 4    | 3          | 4          |
| Sneakers | 4    | 3          | 5          |
| Sneakers | 4    | 3          | 6          |
| T-Shirt  | 7    | 4          | 7          |

## RUNNING TOTAL

A running total is the cumulative sum of a given value and all preceding values in a column.

```
SELECT date, amount,
  SUM(amount) OVER(ORDER BY date)
  AS running_total
FROM sales;
```

## MOVING AVERAGE

A moving average (*a.k.a.* rolling average, running average) is a technique for analyzing trends in time series data. It is the average of the current value and a specified number of preceding values.

```
SELECT date, price,
  AVG(price) OVER(
    ORDER BY date
    ROWS BETWEEN 2 PRECEDING
    AND CURRENT ROW
  ) AS moving_avege
FROM stock_prices;
```

## DIFFERENCE BETWEEN TWO ROWS (DELTA)

```
SELECT year, revenue,
  LAG(revenue) OVER(ORDER BY year)
  AS revenue_prev_year,
  revenue -
    LAG(revenue) OVER(ORDER BY year)
  AS yoy_difference
FROM yearly_metrics;
```

Learn about SQL window functions in our interactive [Window Functions](#) course.